

SHUBHAM PANDEY

Portfolio | LinkedIn | GitHub | Email | 8218554651

Summary

AI/ML Engineer and final-year B.Tech student specializing in Generative AI, Agentic AI, Machine Learning, Deep Learning, and Computer Vision. Experienced in building enterprise AI applications using Large Language Models (LLMs), Multi-Agent Systems, Retrieval-Augmented Generation (RAG), LangChain, LangGraph, FastAPI, and Vector Databases. Strong foundation in Data Structures & Algorithms (DSA), Object-Oriented Programming (OOP), System Design, and Software Engineering with hands-on experience developing scalable, production-ready AI solutions. Seeking AI/ML, Generative AI, or Software Engineering roles to build intelligent, high-impact products.

Education

Bennett University	2023 - 2027
Bachelor of Technology (B.Tech) in Computer Science Engineering	CGPA: 8.0 / 10.0
Shiksha Bharati Senior Secondary School (CBSE)	2022 - 2023
Class XII	84% / 100
Shiksha Bharati Senior Secondary School (CBSE)	2020 - 2021
Class X	91% / 100

Skills

AI/ML & Agentic Frameworks: LangChain, LangGraph, CrewAI, RAG Pipelines, Prompt Engineering, Multi-Agent Systems, LLM Integration, Groq API, ChromaDB, FAISS, Vector Databases
Deep Learning & Machine Learning: PyTorch, TensorFlow, Scikit-Learn, Neural Networks, NLP, Transfer Learning, Object Detection (YOLO), Computer Vision, GANs, Model Fine-Tuning
Data & Analytics: Pandas, NumPy, TF-IDF, Cosine Similarity, Spline Interpolation, Data Preprocessing, Feature Engineering
Programming Languages: Python, C++, JavaScript, TypeScript
Web Frameworks & Deployment: FastAPI, Flask, React, Vite, Tailwind CSS, Streamlit, Uvicorn, REST APIs, Asynchronous Programming (AsyncIO, HTTPX)
Databases, DevOps & Core CS: MongoDB, SQL, Git, GitHub, Docker, CI/CD, Data Structures & Algorithms (DSA), Object-Oriented Programming (OOP), System Design

Experience

BCG X (Boston Consulting Group) — GenAI Job Simulation	View Certificate June 2026
<i>Generative AI Engineer</i>	
– Built a corporate financial chatbot by ingesting and parsing unstructured regulatory disclosures using Pandas, reducing document review time by 30% for enterprise stakeholders – Extracted and structured critical business performance metrics, operational risk signals, and KPI analytics from raw filings, improving data accessibility by 25% for downstream AI reasoning tasks – Engineered context-aware prompt workflows and modular system prompts for enterprise language models, boosting system alignment metrics by 15%	
JPMorgan Chase & Co. — Quantitative Research Job Simulation	View Certificate June 2026
<i>Quantitative Research Analyst</i>	
– Developed a custom cubic spline interpolation model in Python to forecast commodity prices across a 12-month forward horizon, applying quantitative finance and time-series modeling techniques – Designed a reusable asset valuation framework to dynamically estimate baseline asset configurations, reducing manual evaluation overhead by 20% and enhancing contractual profitability metrics – Optimized financial risk evaluation pipelines by implementing programmatic borrower segmentation using credit score bucketing, accelerating risk assessment workflows by 15%	

Projects

AI Technical Interview Coach

GitHub | July 2026

Core Frameworks: Python, React, TypeScript, FastAPI, LangGraph, OpenAI API

- Built a resume-driven mock interview simulator that generates role-specific technical and behavioral questions from uploaded PDFs or Word resumes, eliminating generic questioning.
- Orchestrated adaptive interview flow with LangGraph multi-agent routing, dynamically triggering follow-up probes when responses score below 65/100.
- Delivered multimodal feedback using Whisper, OpenCV, and MediaPipe to evaluate speech clarity, filler words, eye contact, and posture, then surfaced personalized learning paths.

Finvexis AI:Enterprise Multi-Agent Platform

GitHub | April 2026

Core Frameworks: Python, React, TypeScript, Web Services, Multi-Agent Systems, System Design

- Architected a modular multi-agent orchestration system integrating 4 enterprise domains (Business, Finance, Sales, HR) via state machines, accelerating cross-departmental data synthesis by 30%
- Implemented dedicated asynchronous microservices to isolate business logic per domain and handle 500+ concurrent operational requests
- Delivered an executive-facing React dashboard with real-time analytics, structured financial forecasting tools, and budget modules, improving reporting efficiency by 25%

PolyDoc-Chat:Document Intelligence Engine

GitHub | February 2026

Core Frameworks: Python, React, Vector Databases, Cloud APIs, Large Language Models, Semantic Search

- Built a Retrieval-Augmented Generation system supporting 6 distinct file types including digital documents and spreadsheet formats, serving as an enterprise question-answering engine
- Optimized database search and retrieval performance with language models, achieving an end-to-end latency of 0.69 seconds and a retrieval latency of 16.6 milliseconds, a 40%+ improvement over baseline configurations
- Implemented local embedding generation with structural chunk mapping to guarantee 100% citation-backed, hallucination-resistant responses

Co-Drive:Traffic Sign Recognition System

GitHub | January 2026

Core Frameworks: Python, React, Computer Vision, Deep Learning, Deep Neural Networks (DNN)

- Trained and fine-tuned a custom object detection model achieving a verified 99.05% accuracy on sign recognition patterns
- Constructed hardware-accelerated inference pipelines processing image, video, and live webcam streams with a latency of 19 to 25 milliseconds, enabling real-time edge-compatible deployment
- Integrated browser-native high-performance video output via advanced compression encoding with interactive interfaces, improving processing efficiency by 30%

MoodFlix:Conversational Discovery Application

GitHub | March 2026

Core Frameworks: Python, React, Machine Learning, Natural Language Processing, Parallel Pipelines

- Developed a semantic discovery engine using vectorization, cosine similarity, and tokenization to map natural language query intents to relevant database metadata
- Programmed real-time metadata enrichment by implementing asynchronous parallel pipelines, reducing query latency by 35% in live workflows
- Integrated a conversational layer powered by cloud language models alongside fuzzy matching algorithms for typo-tolerance, increasing user engagement metrics by 20%

Certifications

Apply Generative Adversarial Networks (GANs) — DeepLearning.AI	View Certificate
Build Basic Generative Adversarial Networks (GANs) — DeepLearning.AI	View Certificate
Neural Networks and Deep Learning — DeepLearning.AI	View Certificate
Machine Learning: Classification — University of Washington	View Certificate
Statistics for Data Science with Python — IBM	View Certificate